

Lecture 2: Direct Sampling Methods

Deniz Akyildiz

MATH60047/70047 – Stochastic Simulation

October 11, 2022

**Imperial College
London**

Background

Pseudo-random number generation

Uniform random number generation

Direct sampling for other distributions

We will denote densities here with $p(\cdot)$. For generic densities, we use the notation $X \sim p(x)$, **not** $p_X(x)$ (unless necessary).

In this context, $Y \sim p(y)$ will denote *another density*. Sometimes we will distinguish $p_X(x)$ and $p_Y(y)$ (especially when we do transformation of r.v.s).

Why? When we go into Bayesian inference, this makes it slightly tedious to write things down.

Recall that $p(x)$ is a *function* that integrates to one:

$$\int p(x) dx = 1.$$

A probability measure is when we write \mathbb{P} . We won't need the notion throughout the course.

But let's clarify one thing:


$$\mathbb{P}(x_1 \leq X \leq x_2) = \mathbb{P}(X \in [x_1, x_2]) = \int_{x_1}^{x_2} p(x) dx.$$

$\mathbb{P}(\cdot)$ *measures* sets (intervals etc.) using the density.

What is the probability of a point for a continuous variable X ?


$$\mathbb{P}(X = c) = ?$$

Is it $p(c)$?



$$p(0) = 1.2615$$

$p(0)$ is **not** the probability of $X = 0$. For continuous variables, probability only makes sense on intervals (or sets in higher dimensions).



$$\mathbb{P}(-0.1 \leq X \leq 0.1) = \int_{-0.1}^{0.1} p(x) dx = 0.248$$

So pointwise evaluations *could be* bigger than 1 in continuous case.
Don't get confused!

For discrete sets, we denote the state-space $S = \{s_1, \dots, s_K\}$ for K possibilities.

► $K = 6$ for a die and $s_1 = 1, \dots, s_6 = 6$.

Probability mass function:

$$p(s_k) \in [0, 1]$$

and $\sum_{k=1}^K p(s_k) = 1$.

Discrete case is easy:

$$\mathbb{P}(X = s_k) = p(s_k).$$

The CDF is defined as (for continuous variables)

$$F_X(x) = \mathbb{P}(X \leq x) = \int_{-\infty}^x p(x') dx'.$$

or for discrete variables

$$F(k) = \sum_{i=-\infty}^k p(s_i)$$

Definition 1

A sequence of pseudo-random numbers u_1, \dots, u_n is a deterministic sequence of numbers with *good enough* statistical properties that match the distribution we want to simulate from.


Good, but how do we test “statistical properties”?

Say you are given a sequence of numbers x (an `np.array`).

- ▶ Check moments: compute the mean (`np.mean(x)`) or the variance (`np.var(x)`) of the sample and check if they match the parameters
- ▶ Check histogram against the density (if you know it)
 - ▶ `plt.hist(x, bins=100, density=True)`

More on this later, but for now these will be enough.

What are pseudo-random numbers?



$$X^{(i)} \sim p(x)$$

What are pseudo-random numbers?

Why do we need them?

It is (literally) impossible to generate genuinely random numbers on computers.

- ▶ You can flip a coin every time you need a binary number
 - ▶ Is it really unbiased though?¹
- ▶ Throw a die

What other things can give you a truly random number?

- ▶ You can use www.random.org
- ▶ On a computer
 - ▶ Try to measure some inner thermal noise (of circuits)
 - ▶ Measure atmospheric noise

As you can see, these are not very practical.

¹Diaconis, P., Holmes, S., & Montgomery, R. (2007). Dynamical bias in the coin toss. *SIAM review*, 49(2), 211-235.

What are pseudo-random numbers?

Why do we need them?

If we want to simulate randomness, we need to obtain a way that is

- ▶ Repeatable
- ▶ Cheap

It has become an entire research topic to design *deterministic* algorithms which gives samples that match the desired characteristics.

We will start from the simplest: The uniform distribution.

Uniform pseudo-random numbers

The technology that underlies modern civilization

The key to simulate many (many) other random variables is to be able to simulate uniform random numbers.

We denote the task

$$U \sim \text{Unif}(u; 0, 1).$$

More precisely

$$U \sim p(u) = 1 \quad \text{for } 0 \leq u \leq 1.$$

We will look into the standard way of doing it:

- ▶ Linear congruential random number generators

These methods are based on generating a *deterministic linear recursion* with a careful design.

Linear congruential generators (LCGs from now on) are based on simulating a recursion:

$$x_{n+1} \equiv ax_n + b \pmod{m}$$

where x_0 is the **seed**, m is the **modulus**, b is the **shift**, and a is the **multiplier**.

- ▶ m is an integer
- ▶ $x_0, a, b \in \{0, \dots, m-1\}$.

Given $x_n \in \{0, \dots, m-1\}$, we generate the uniform random numbers

$$u_n = \frac{x_n}{m} \in [0, 1) \quad \forall n.$$

Example code (try and make it work!)

```
import numpy as np
import matplotlib.pyplot as plt

def lcg(a, b, m, n, x0):
    x = np.zeros(n)
    u = np.zeros(n)
    x[0] = x0
    u[0] = x0 / m
    for k in range(1, n):
        x[k] = (a * x[k - 1] + b) % m
        u[k] = x[k] / m
    return u
```

A few things to know about LCGs:

- ▶ They generate *periodic* sequences.




Figure: $m = 2048$, $a = 43$, $b = 0$, $x_0 = 1$.

period $T \leq m$ (m : the modulus).

- ▶ Full period: $T = m$

Choice of good parameters rely on some theory, some art.

Uniform pseudo-random numbers

The technology that underlies modern civilization

Wikipedia has a list of parameters for professional implementations:



Parameters in common use [\[edit\]](#)

The following table lists the parameters of LCGs in common use, including built-in `rand()` functions in [runtime libraries](#) of various [compilers](#). This table is to show popularity, not examples to emulate; many of these parameters are poor. Tables of good parameters are available.^{[10][2]}

Source	modulus <i>m</i>	multiplier <i>a</i>	increment <i>c</i>	output bits of seed in <code>rand()</code> or <code>Random(L)</code>
Zx81	$2^{16} + 1$	75	74	
Numerical Recipes from the "quick and dirty generators" list, Chapter 7.1, Eq. 7.1.6 parameters from Knuth and H. W. Lewis	2^{32}	1664525	1013904223	
Borland C/C++	2^{32}	22695477	1	bits 30..16 in <code>rand()</code> , 30..0 in <code>lrand()</code>
glibc (used by GCC) ^[17]	2^{31}	1103515245	12345	bits 30..0
ANSI C: Watcom, Digital Mars, CodeWarrior, IBM VisualAge C/C++ ^[18] C90, C99, C11: Suggestion in the ISO/IEC 9899, ^[19] C17	2^{31}	1103515245	12345	bits 30..16
Borland Delphi, Virtual Pascal	2^{32}	134775813	1	bits 63..32 of (seed \times L)
Turbo Pascal	2^{32}	134775813 (8088405 ₁₆)	1	
Microsoft Visual/Quick C/C++	2^{32}	214013 (343FD ₁₆)	2531011 (269EC3 ₁₆)	bits 30..16
Microsoft Visual Basic (6 and earlier) ^[20]	2^{24}	1140671485 (43FD43FD ₁₆)	12820163 (C39EC3 ₁₆)	
RtlUniform from Native API ^[21]	$2^{31} - 1$	2147483629 (7FFFFFFD ₁₆)	2147483587 (7FFFFFFC3 ₁₆)	
Apple CarbonLib, C++11's				

from: https://en.wikipedia.org/wiki/Linear_congruential_generator


Why professional implementation? Consider this Gaussian:




Uniform pseudo-random numbers

The technology that underlies modern civilization

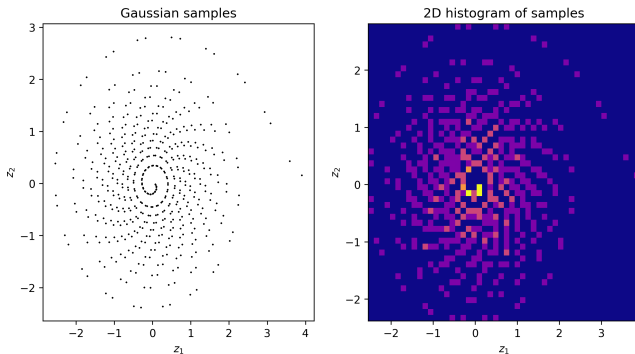
Samples drawn using Python's uniform




Histogram of samples



Samples drawn using the random number generator I showed before:



Samples drawn using the second sampler on Wiki:



Going forward, we will mostly assume that we will have access to a uniform random number generator.

- ▶ In assignments, we won't ask you to implement LCG.
- ▶ When asked for $U \sim \text{Unif}(0, 1)$, you can instead use

```
np.random.uniform(0, 1, n)
```

where `n` is the number of samples you want to draw.

Next up: Sampling from general $p(x)$

Exact sampling of distributions

How to simulate from any $p(x)$?

Simulating from a given $p(x)$ is an endless research area (simulation, sampling, generative models) and still flourishing.

We will start by describing some general methods to sample from more general distributions.

This lecture (direct sampling):

- ▶ Inversion
- ▶ Transformation

Direct sampling of distributions

Inversion

Imperial College
London

THE INSTITUTE FOR ADVANCED STUDY
SCHOOL OF MATHEMATICS
PRINCETON, NEW JERSEY

May 21, 1947

Mr. Stan Ulam
Post Office Box 1663
Santa Fe
New Mexico

Dear Stan:

Thanks for your letter of the 19th. I need not tell you that Klari and I are looking forward to the trip and visit at Los Alamos this Summer. I have already received the necessary papers from Carson Mark. I filled out and returned mine yesterday; Klari's will follow today.

I am very glad that preparations for the random numbers work are to begin soon. In this connection, I would like to mention this: Assume that you have several random number distributions, each equidistributed in $0, 1 : (x^i), (y^i), (z^i), \dots$. Assume that you want one with the distribution function (density) $f(\xi) d\xi : (\xi^i)$. One way to form it is to form the cumulative distribution function: $g(\xi) = \int_0^\xi f(\xi) d\xi$ to invert it $h(x) = \xi \Leftrightarrow x = g(\xi)$, and to form $\xi^i = h(x^i)$ with this $h(x)$, or some approximant polynomial. This is, as I see, the method that you have in mind.

The inversion technique is based on the following theorem:

Theorem 2

Consider a random variable X with a CDF F_X . Then the random variable

$$Y = F_X(X)$$

is uniformly distributed.

Proof. Consider any continuous random variable X , we define $Y = F_X(X)$. For $y \in [0, 1]$,

$$\begin{aligned}F_Y(y) &= \mathbb{P}(Y \leq y) \\&= \mathbb{P}(F_X(X) \leq y) \\&= \mathbb{P}(X \leq F_X^{-1}(y)) && \text{Why?} \\&= F_X(F_X^{-1}(y)) \\&= y,\end{aligned}$$

which is the CDF of the standard uniform distribution.

Note that above result is written for the case where F_X^{-1} exists, i.e., the CDF is continuous. If this is not the case, one can define the generalised inverse function,

$$F_X^-(u) = \min\{x : F_X(x) \geq u\}.$$

Going back to statement: If $X \sim p(x)$ with $F_X(x) = \int_{-\infty}^x p(x')dx'$, we know that

$$Y = F_X(X)$$

is uniform. Then this suggests *inverting* this process:

- ▶ Sample $U \sim \text{Unif}([0, 1])$,
- ▶ Draw $X = F_X^{-1}(U)$.

Of course, this is limited to the cases where we *can* invert the CDF.

Exact sampling of distributions

Inversion: Discrete (categorical) distribution

Let us consider some examples.

The most generic one is the **discrete (categorical) distribution**. For $K \geq 1$ (integer), define K states s_1, \dots, s_K where

$$p(s_k) \in [0, 1] \quad \text{where} \quad \sum_{k=1}^K p(s_k) = 1.$$

Simpler than it looks, consider the die:

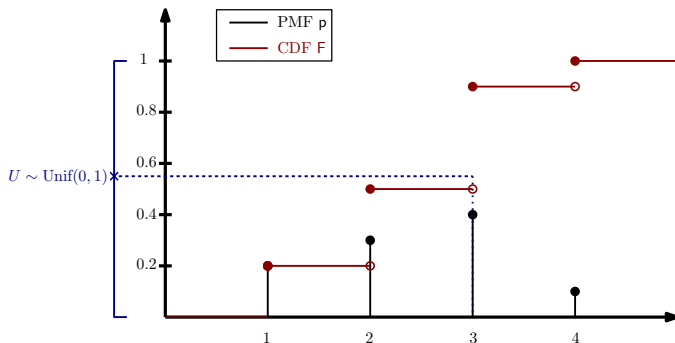
$$s_k = k \text{ (the face of die)}$$

and their probabilities

$$p(s_k) = 1/6.$$



How does the sampling work?



- ▶ Draw $U \sim \text{Unif}(0, 1)$
- ▶ Choose $F_X^-(u) = \min\{x : F_X(x) \geq u\}$ generic for discrete dist.

Exact sampling of distributions

Inversion: Discrete (categorical) distribution

Some starters:

```
import numpy as np
import matplotlib.pyplot as plt

w = np.array([0.2, 0.3, 0.4, 0.1]) # pmf
s = np.array([1, 2, 3, 4])        # support (states)

def discrete_cdf(w):
    return np.cumsum(w)

cw = discrete_cdf(w)

def plot_discrete_cdf(w, cw):
    fig, ax = plt.subplots(1, 2, figsize=(20, 5))
    ax[0].stem(s, w)
    ax[1].plot(s, cw, 'o-', drawstyle='steps-post')
    plt.show()

plot_discrete_cdf(w, cw)
```


Exact sampling of distributions

Inversion: Exponential distribution

The exponential density

$$p(x) = \text{Exp}(x; \lambda) = \lambda e^{-\lambda x}.$$

for $x \geq 0$. Otherwise $p(x) = 0$.



Exact sampling of distributions

Inversion: Exponential distribution

$$p(x) = \text{Exp}(x; \lambda) = \lambda e^{-\lambda x}.$$

We calculate the CDF

$$\begin{aligned} F_X(x) &= \int_0^x p(x') dx', \\ &= \lambda \int_0^x e^{-\lambda x'} dx', \\ &= \lambda \left[-\frac{1}{\lambda} e^{-\lambda x'} \right]_{x'=0}^x \\ &= 1 - e^{-\lambda x}. \end{aligned}$$

Calculate the reverse?

Exact sampling of distributions

Inversion: Exponential distribution

Deriving the inverse:

$$\begin{aligned}u &= 1 - e^{-\lambda x} \\ \implies x &= -\frac{1}{\lambda} \log(1 - u) \\ \implies F_X^{-1}(u) &= -\lambda^{-1} \log(1 - u).\end{aligned}$$

So what is the **algorithm**?

- ▶ Generate $u_i \sim \text{Unif}([0, 1])$
- ▶ $x_i = -\lambda^{-1} \log(1 - u_i)$.

Implement this (exercises).

Consider the Cauchy distribution

$$p(x) = \frac{1}{\pi(1+x^2)}.$$

Given

$$p(x) = \frac{1}{\pi(1+x^2)},$$

the CDF is given by:

$$F_X(x) = \int_{-\infty}^x p(x')dx.$$

How do we compute this integral? (at the end, if time permits)

We obtain:

$$F_X(x) = \frac{1}{2} + \pi^{-1} \tan^{-1} x.$$

The inverse is straightforward:

$$F_X^{-1}(u) = \tan \left[\pi \left(u - \frac{1}{2} \right) \right].$$

Another one: Given

$$p(k) = \text{Pois}(k; \lambda) = \frac{\lambda^k e^{-\lambda}}{k!}.$$

Design a sampling scheme using inversion.

The CDF is given:

$$F(k) = \sum_{i=1}^k \text{Pois}(k; \lambda) = e^{-\lambda} \sum_{i=0}^k \frac{\lambda^i}{i!}.$$

How do you use inversion? Recall the discrete example.

Wait until it hits:

▶ Draw $U \sim \text{Unif}([0, 1])$



$$P = \min \left\{ k \in \mathbb{N} : U \leq e^{-\lambda} \sum_{i=0}^k \frac{\lambda^i}{i!} \right\}.$$

This is done via a `while` statement.

Exact sampling of distributions

Inversion: Is Gaussian possible?

Let $p(x) = \mathcal{N}(x; \mu, \sigma^2)$. Can we use inversion?

No. F_X^{-1} is impossible to compute and hard to approximate.

Exact sampling of distributions

Transformation method

Inversion is a special case of a general method called *transformation method*.

Transformation method:

- ▶ Sample $U_i \sim \text{Unif}(u; 0, 1)$
- ▶ Transform: $X_i = g(U_i)$.

Inversion is just setting $g = F_X^{-1}$.

Exact sampling of distributions

Transformation method: Sampling a custom uniform

The simplest example can be seen from sampling a uniform on $[a, b]$ using a uniform on $[0, 1]$.

- ▶ Draw $U_i \sim \text{Unif}(u; 0, 1)$
- ▶ Set $X_i = g(U_i) = (b - a)U_i + a$

then $X_i \sim \text{Unif}(x; a, b)$.

For general g , how do we compute the density?

If $X \sim p_X(x)$ and $Y = g(X)$, what is $p_Y(y)$?

$$p_Y(y) = p_X(g^{-1}(y)) |J_{g^{-1}}(y)|$$

where J is the Jacobian of the inverse mapping g^{-1} , evaluated at y .

Exact sampling of distributions

Transformation method: Sampling a Gaussian (Box and Müller, 1958)

Theorem 3 (Box-Müller method)

Let X_1, X_2 be independent r.v.'s respectively where

$$X_1 \sim \text{Exp}\left(\frac{1}{2}\right),$$
$$X_2 \sim \text{Unif}(0, 2\pi).$$

Then $Y_1 = \sqrt{X_1} \cos X_2$ and $Y_2 = \sqrt{X_1} \sin X_2$ are independent and $\mathcal{N}(0, 1)$ -distributed.

A transformation method with

$$(y_1, y_2) = g(x_1, x_2) = (\sqrt{x_1} \cos x_2, \sqrt{x_1} \sin x_2).$$

Exact sampling of distributions

Transformation method: Sampling a Gaussian (Box and Müller, 1958)

Proof. How to compute the density $p(y_1, y_2)$? Use the transformation of random variables

$$p_{y_1, y_2}(y_1, y_2) = p_{x_1, x_2}(g^{-1}(y_1, y_2)) |J_{g^{-1}}(y_1, y_2)|$$

where $J_{g^{-1}}$ is the Jacobian of the inverse. What is g^{-1} ?

Recall $(y_1, y_2) = g(x_1, x_2) = (\sqrt{x_1} \cos x_2, \sqrt{x_1} \sin x_2)$. We have

$$x_1 = y_1^2 + y_2^2, \quad \text{as } \cos^2 + \sin^2 = 1.$$

and

$$\frac{\sin x_2}{\cos x_2} = \frac{y_2}{y_1}$$

which leads to

$$x_2 = \arctan(y_2/y_1).$$

Exact sampling of distributions

Transformation method: Sampling a Gaussian (Box and Müller, 1958)

Therefore, $g^{-1} : \mathbb{R}^2 \rightarrow \mathbb{R}^2$

$$g^{-1}(y_1, y_2) = (g_1^{-1}, g_2^{-1}) = (y_1^2 + y_2^2, \arctan(y_2/y_1)) .$$

Compute the Jacobian

$$\begin{aligned} J_{g^{-1}} &= \begin{bmatrix} \partial g_1^{-1} / \partial y_1 & \partial g_1^{-1} / \partial y_2 \\ \partial g_2^{-1} / \partial y_1 & \partial g_2^{-1} / \partial y_2 \end{bmatrix} \\ &= \begin{bmatrix} 2y_1 & 2y_2 \\ \frac{1}{1+(y_2/y_1)^2} \frac{-y_2}{y_1^2} & \frac{1}{1+(y_2/y_1)^2} \frac{1}{y_1} \end{bmatrix} \end{aligned}$$

Hence, the determinant is:

$$|J_{g^{-1}}| = 2.$$

Exact sampling of distributions

Transformation method: Sampling a Gaussian (Box and Müller, 1958)

Remember...

$$p_{y_1, y_2}(y_1, y_2) = p_{x_1, x_2}(g^{-1}(y_1, y_2)) |J_{g^{-1}}(y_1, y_2)|.$$

and

$$g^{-1}(y_1, y_2) = (g_1^{-1}, g_2^{-1}) = (y_1^2 + y_2^2, \arctan(y_2/y_1)).$$

Let's write it out

$$\begin{aligned} p_{y_1, y_2}(y_1, y_2) &= \text{Exp}(g_1^{-1}; 1/2) \text{Unif}(g_2^{-1}; 0, 2\pi) |J_{g^{-1}}| \\ &= \frac{1}{2} e^{-\frac{1}{2}(y_1^2 + y_2^2)} \frac{1}{2\pi} 2 \\ &= \mathcal{N}(y_1; 0, 1) \mathcal{N}(y_2; 0, 1). \end{aligned}$$



Exact sampling of distributions

Transformation method: Sampling a Gaussian (Box and Müller, 1958)

```
import matplotlib.pyplot as plt
import numpy as np

def box_muller(n):

    # Your code



    return y_1, y_2

n = 100000
y_1, y_2 = box_muller(n)

fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].scatter(y_1, y_2, s=0.1, c='k')
axes[1].hist2d(y_1, y_2, bins=50, cmap='plasma')
plt.show()
```


Exact sampling of distributions

Transformation method: Sampling a Gaussian (Box and Müller, 1958)



Exact sampling of distributions

Transformation method: A puzzle

Define the problem. We draw

$$r \sim \text{Unif}(0, 1),$$

$$\theta \sim \text{Unif}(0, 2\pi).$$

How to get uniform points on a circle?

Let's have a little poll.

Exact sampling of distributions

Transformation method: A puzzle

Define the problem. We draw

$$r \sim \text{Unif}(0, 1),$$

$$\theta \sim \text{Unif}(0, 2\pi).$$

How to get uniform points on a circle? Which one of the following?


▶ $x_1 = \sqrt{r} \cos \theta$ and $x_2 = \sqrt{r} \sin \theta$

▶ $x_1 = r \cos \theta$ and $x_2 = r \sin \theta$

▶ $x_1 = r^2 \cos \theta$ and $x_2 = r^2 \sin \theta$


Exact sampling of distributions

Transformation method: A puzzle




$$y_1 = \sqrt{r} \sin(\theta)$$

$$x_1 = \sqrt{r} \cos(\theta)$$



$$y_2 = r \sin(\theta)$$

$$x_2 = r \cos(\theta)$$



$$y_3 = r^2 \sin(\theta)$$

$$x_3 = r^2 \cos(\theta)$$

Exact sampling of distributions

Transformation method: A puzzle

The transformation is *exactly* the same as Box-Müller (exercise).

It is possible to prove that

$$p_{x_1, x_2}(x_1, x_2) = \begin{cases} \frac{1}{\pi} & \text{if } x_1^2 + x_2^2 < 1 \\ 0 & \text{otherwise,} \end{cases}$$

which is the uniform distribution on the circle.

Exact sampling of distributions

Transformation method: A puzzle

In class exercise: Prove this result.

$$\begin{aligned} p_{x_1, x_2}(x_1, x_2) &= \text{Unif}(g_1^{-1}; 0, 1) \text{Unif}(g_2^{-1}; 0, 2\pi) |J_{g^{-1}}| \\ &= \frac{1}{2\pi} 2 \quad \text{for } x_1^2 + x_2^2 < 1, \\ &= \frac{1}{\pi} \quad \text{for } x_1^2 + x_2^2 < 1. \end{aligned}$$

If $X \sim \mathcal{N}(0, 1)$, derive the distribution of

$$Y = \sigma X + \mu.$$

Exact sampling of distributions

Transformation method: Another exercise

The inverse transform is:

$$g^{-1}(y) = \frac{y - \mu}{\sigma}.$$

Therefore,

$$p_Y(y) = p_X(g^{-1}(y)) \left| \frac{dg^{-1}}{dy} \right|,$$

which is

$$p_Y(y) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right) \frac{1}{\sigma} = \mathcal{N}(\mu, \sigma^2)$$

method that you have in mind. ~~polynomial~~. This is, as I see, the

An alternative, which works if ξ and all values of $f(\xi)$ lie in 0, 1, is this: Scan pairs x^i, y^i and use or reject x^i, y^i according to whether $y^i \leq f(x^i)$ or not. In the first case, put $\xi^d = x^i$ in the second case form no ξ^d at that step.

The second method may occasionally be better than

Rejection sampling

- 📎 Box, GEP and Mervin E Müller (1958). “A Note on the Generation of Random Normal Deviates”. In: *The Annals of Mathematical Statistics* 29.2, pp. 610–611.